FREE GUIDE

dabeer.ai

The Super Agent Guide

What a super agent is, how DeepSeek built one, and how to try it yourself. In plain words.

Based on DeepSeek's open-source Super-Agent

Released Oct 31, 2026 under the MIT licence (repo: deepseek-ai/dsh-superagent-prototype). Everything here is checked against DeepSeek's own documentation.

By Dabeer Shaikh · November 2026

dabeer.ai

What's inside

Part 1 · The idea
Chatbot vs agent vs super agent · The one rule behind it

Part 2 · How it works
The flow · 9 building blocks · Router and Workers · Who can give it work · The life of a task · The panel

Part 3 · Try it yourself
What you need · Install · Your first request

Part 4 · Using it well
Good first jobs · Rules for the operator · Honest limits · What's next

How to read this

Each page builds on the one before it, the same way DeepSeek wrote its own docs. No page needs a word that hasn't been explained yet.

dabeer.ai

Chatbot, agent, super agent

	💬 Chatbot	🤖 Agent	🦸 Super agent
Who starts it?	You	You	Your team's request or a team rule
What does it do?	Answers	Does one task	Runs many tasks at once
Where does it live?	A chat window	Your laptop / IDE	Your team chat + GitHub
Who watches it?	You, every reply	You, every step	Nobody, until it needs a person

The analogy

A chatbot is a helpline. An agent is a smart intern at your desk: you give it a task and watch it work. A super agent is a team member in your group chat: someone says "can you fix the login bug?", it picks it up, and comes back with a pull request.

👉 **With an agent, you manage the AI. With a super agent, the AI manages the work.**

dabeer.ai

The one rule behind it

> **"Take as little human attention as possible, and give as much value as possible."**
>
> DeepSeek, 000 Philosophy

Every message, notification, question and setting costs someone's attention. So DeepSeek built the agent to:

🙋 Act only when asked

Or when a rule the team wrote allows it. Work nobody asked for starts in a mode that only records what it would have done.

🤫 Speak only when it matters

It confirms once, stays quiet while working, and writes again only when a result is ready, it needs a decision, or it can't finish.

❓ Ask one question at a time

A short, concrete question instead of a long report.

🛑 Stop by impact, not effort

A small change goes all the way to a pull request. A change that needs a design decision stops with a clear question.

dabeer.ai

The flow, step by step

Example: someone writes `@agent fix the login bug` on a GitHub issue.

1. 📡 **Source**
   GitHub (or Feishu chat) is watched for new events
2. ⚡ **Trigger**
   The new comment becomes one event
3. 🧵 **Thread**
   It joins the issue's conversation history
4. ⚖️ **Eval**
   The Router judges: is the person allowed? Is this real work?
5. 🎫 **Ticket + Mandate**
   A unit of approved work, with a written note of what it may do
6. 🛠️ **Executor**
   A Worker session does the job in its own folder
7. 📬 **Channel**
   The answer goes back to the same issue, often as a pull request

dabeer.ai

9 building blocks, in plain words

Source	Where requests come from: GitHub, or a Feishu chat. Like the office inbox.
Trigger	One event, like a comment. Seen twice, it still makes only one task.
Thread	A conversation in order, like one issue or one chat thread.
Eval	The judgement: start new work, continue it, or stop it.
Mandate	The permission slip: why this work is allowed and how far it may go. Can't be widened later; bigger scope means a new task.
Ticket	One unit of approved work. Only Eval can create one.
Executor	The worker doing the job, in its own session and folder.
Channel	The way back: replies, pull requests, review comments.
Subscription	Keeps a task "listening" to its thread, so later comments reach the worker without anyone tagging the agent again.

dabeer.ai

Router and Workers: the manager and the team

🧭 Router (main session)

Reads every incoming request and decides what is real work.

Has a deliberately **small set of tools**. Shell commands and waits are capped at 5 seconds.

Never writes the code itself. It only routes.

🛠️ Worker (one per task)

Each task gets its own session, named `worker—<ticket id>`, and its own folder.

Long commands move to the background after 5 seconds, so it can still read new messages.

Background work gets up to 30 minutes after each turn.

Why split it this way?

Like a restaurant: the host seats people and takes orders; the cooks each work at their own station. The host never cooks, and one cook's mess never spills into another's dish. That's why many tasks can run in parallel without stepping on each other.

dabeer.ai

Who can give it work?

✅ Explicit request

An allowed person mentions the agent (@agent ...) or messages it directly. This works in every mode.

🤖 Bots can only suggest

A bot on the approved list can create a "candidate", never a direct order. Other bots are ignored.

📜 Proactive work needs a written rule

The team writes rules like "fix failing CI on main". Work nobody asked for only starts if a rule matches.

Accept mode	What happens to proactive work
off	Not considered at all
observe (default)	Judged and recorded as "would accept", but nothing starts
auto	Starts when a rule matches

DeepSeek's advice: start in observe, read what it would have done, then switch to auto.

dabeer.ai

The life of a task

Every open Ticket sits in one stage. The panel shows the label on the right.

Stage	Means	Panel label
Queued	Waiting for a free worker	Queued
Working	A worker is on it	Working
Needs a person	It asked a question	Awaiting answer
Needs a person	Result is ready	Awaiting review
Needs a person	It got stuck or failed	Blocked · Failed
Idle	Done for now, listening for the next comment	Waiting for a Trigger

The nice part

A finished task keeps listening to its thread. If someone comments "also handle the logout page" on the pull request, the same worker picks it back up under the same permission slip. No new tagging needed.

dabeer.ai

The panel: 3 questions

Open "Superagent" in the DeepSeek Harness web sidebar. It answers, in this order:

1. Does anything need me now?
Tasks waiting for an answer or a review come first. A healthy system takes no space.

2. Is it healthy and doing the right things?
Source problems, activity in the last 24 hours, how busy the workers are.

3. What happened to one piece of work?
The full trail: Trigger → Eval decision and reason → Ticket → result.

Tabs: Overview · Source · Trigger · Eval · Execution · Settings.

dabeer.ai

What you need

💻 Tools on your machine

Node.js 22.19+ or 24+, pnpm, git 2.32+, Python 3.8+, tmux, and GitHub CLI (gh) 2.48+.

🐙 A GitHub repo + a separate GitHub account for the agent

The agent comments and opens pull requests as its own user, with write access to the repo. (A GitHub App can't be the agent.) GitHub Enterprise Server isn't supported.

🧠 A model provider

Configured in DeepSeek Harness. The agent uses that profile's default model.

💬 Optional: a Feishu / Lark bot

Only if you want requests from chat, not just GitHub.

⚠️ Run it on a spare machine or VM

Workers run commands without asking, as your OS user. DeepSeek recommends a dedicated OS account that holds only the agent's credentials.

dabeer.ai

Install in 5 commands

From DeepSeek's deployment guide (006). Run once, on a machine with no "superagent" profile yet:

```
npm install -g @deepseek-ai/dsh@0.2.1-rc.3
git clone https://github.com/deepseek-ai/dsh-superagent-prototype dsh-superagent
cd dsh-superagent && pnpm install && pnpm build
dsh superagent --from-default-profile web --dump-config > /dev/null
dsh plugin --profile superagent add "link:$PWD"
```

Easier route

DeepSeek's README suggests letting a coding agent do it with you. Ask it: *"Read SETUP.md in this repository and set up dsh-superagent with me."* It asks one question at a time, shows the plan, then installs.

dabeer.ai

Your first request

1. **Open an issue**
   From an allowed GitHub account, write: @<agent login> add a README section on setup
2. **Wait about a minute**
   It checks GitHub once a minute by default. A Ticket appears on the Overview tab
3. **It confirms once**
   One short reply saying what it will do. Then it goes quiet
4. **It works alone**
   In its own session and folder
5. **It replies in the same issue**
   Often with a pull request for you to review

Review it like a teammate's PR. If something's off, just comment on the PR: the same worker picks it up.

dabeer.ai

Good first jobs

✅ Start with

- Small bug fixes from clear issues
- Adding or fixing tests
- Updating docs and READMEs
- "Where is X handled in the code?"
- Investigating a failing build

⏸️ Not yet

- Big design decisions (it should stop and ask)
- Changes across many services
- Anything nobody will review
- Mass pull requests or busywork

DeepSeek's rules for the operator

Start narrow

A short allowed list, and proactive work in observe mode. Widen once it looks right.

Tune for less noise

If it's noisy, make its prompts shorter and stricter, or pause it. Don't add more messages.

Own what it says

It posts under an account your team sees daily. Read it, and fix the prompts when it's wrong.

dabeer.ai

Honest limits

🧪 It's a research prototype

DeepSeek: "not an official DeepSeek product", may contain bugs, not production-ready. The public code differs from their internal version, and some ideas in their design article aren't built yet.

📊 No benchmarks published

DeepSeek hasn't released performance numbers, so treat any "better than X" claim with care.

🤝 Agent teams are still hard

A January study (CooperBench) found coding agents did about 30% worse working as a team than alone. It didn't test DeepSeek's agent, but it shows why the Router/Worker split matters.

🐙 One repo, GitHub.com only

The prototype works on one repository, and not on GitHub Enterprise Server.

dabeer.ai

Where this is going

The big labs are all moving from "an AI you talk to" toward "AI that runs work for a team":

🐋 DeepSeek

Open-sourced Super-Agent (Oct 31) on top of its open-source Harness. Bloomberg reports it's aimed at Claude Code.

🐙 GitHub

At Universe, showed a Copilot app that runs agent work across several repos in parallel sessions.

What it means for you

The skill shifts from "writing the prompt" to **writing good issues, clear rules and good reviews.** The people who can describe work clearly and review it fast will get the most out of super agents.

Found this useful?

Follow Dabeer Shaikh on LinkedIn and @dabeer.ai on Instagram for more plain-words AI guides.

dabeer.ai

Sources

- DeepSeek, dsh-superagent-prototype README and docs 000–006 · github.com/deepseek-ai/dsh-superagent-prototype
- The Next Web, "DeepSeek open-sources Super-Agent, a coding agent that reads team chats", Nov 1, 2026 (citing Bloomberg and SCMP)
- Startup Fortune, "DeepSeek Open Sources the Architecture Behind Its Super-Agent", Nov 1, 2026

Commands and version numbers are from DeepSeek's docs as of Nov 2, 2026. Check the repo for the latest before you install.